

johanan-119-lab5-1

October 25, 2024

NNDL

Lab Program - 5

Submitted by: Johanan Joshua (2347119)

Lab 5 Question:

Topic: Implementing CNN on the Fashion-MNIST Dataset

Objective: In this lab, you will implement a Convolutional Neural Network (CNN) using the Intel Image Classification dataset. Your task is to train a CNN model to classify these images with high accuracy. (Link:- <https://www.kaggle.com/datasets/puneet6060/intel-image-classification>)

Tasks:

1. Dataset Overview:
 - Visualize a few samples from the dataset, displaying their corresponding labels.
2. Model Architecture:
 - Design a CNN model with at least 3 convolutional layers, followed by pooling layers and fully connected (dense) layers.
 - Experiment with different kernel sizes, activation functions (such as ReLU), and pooling strategies (max-pooling or average pooling).
 - Implement batch normalization and dropout techniques to improve the generalization of your model.
3. Model Training:
 - Split the dataset into training and test sets.
 - Compile the model using an appropriate loss function (categorical cross- entropy) and an optimizer (such as Adam or SGD).
 - Train the model for a sufficient number of epochs, monitoring the training and validation accuracy.
4. Evaluation:
 - Evaluate the trained model on the test set and report the accuracy.
 - Plot the training and validation accuracy/loss curves to visualize the model's performance.
 - Display the confusion matrix for the test set to analyze misclassified samples.
5. Optimization :

- Experiment with data augmentation techniques (rotation, flipping, zooming) to further improve the model's performance.
- Fine-tune hyperparameters like learning rate, batch size, and the number of filters in each layer.

Solution:

Step 1: Import Necessary Libraries

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Dropout, BatchNormalization
from sklearn.metrics import confusion_matrix
```

This step involves importing necessary libraries such as TensorFlow, Matplotlib, and Seaborn, which provide functions for building neural networks, visualizing data, and performing statistical analyses.

```
[2]: import kagglehub

# Download latest version
path = kagglehub.dataset_download("puneet6060/intel-image-classification")

print("Path to dataset files:", path)
```

Downloading from

https://www.kaggle.com/api/v1/datasets/download/puneet6060/intel-image-classification?dataset_version_number=2...

100%| | 346M/346M [00:16<00:00, 21.8MB/s]

Extracting files...

Path to dataset files: /root/.cache/kagglehub/datasets/puneet6060/intel-image-classification/versions/2

```
[3]: !pip install kaggle
```

Requirement already satisfied: kaggle in /usr/local/lib/python3.10/dist-packages (1.6.17)

Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.10/dist-packages (from kaggle) (1.16.0)

Requirement already satisfied: certifi>=2023.7.22 in

```

/usr/local/lib/python3.10/dist-packages (from kaggle) (2024.8.30)
Requirement already satisfied: python-dateutil in
/usr/local/lib/python3.10/dist-packages (from kaggle) (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-
packages (from kaggle) (2.32.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages
(from kaggle) (4.66.5)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.10/dist-
packages (from kaggle) (8.0.4)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-
packages (from kaggle) (2.2.3)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages
(from kaggle) (6.1.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-
packages (from bleach->kaggle) (0.5.1)
Requirement already satisfied: text-unidecode>=1.3 in
/usr/local/lib/python3.10/dist-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests->kaggle) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-
packages (from requests->kaggle) (3.10)

```

```
[4]: !kaggle datasets download -d puneet6060/intel-image-classification
```

```

Dataset URL: https://www.kaggle.com/datasets/puneet6060/intel-image-
classification
License(s): copyright-authors
Downloading intel-image-classification.zip to /content
100% 346M/346M [00:16<00:00, 22.7MB/s]
100% 346M/346M [00:16<00:00, 21.7MB/s]

```

```

[5]: import zipfile

# Unzipping with suppressed output
with zipfile.ZipFile("intel-image-classification.zip", "r") as zip_ref:
    zip_ref.extractall("intel-image-classification")

print("Unzipping completed successfully.")

```

Unzipping completed successfully.

```

[6]: # Define the data directory
data_dir = '/root/.cache/kagglehub/datasets/puneet6060/
↳intel-image-classification/versions/2/seg_train/seg_train'


# Class names
class_names = os.listdir(data_dir)

```

```

# Plot a few samples from each class
plt.figure(figsize=(12, 8))
for i, class_name in enumerate(class_names):
    img_path = os.path.join(data_dir, class_name)
    images = os.listdir(img_path)[:5] # First 5 images
    for j, img in enumerate(images):
        img_full_path = os.path.join(img_path, img)
        img_array = plt.imread(img_full_path)
        plt.subplot(len(class_names), 5, i * 5 + j + 1)
        plt.imshow(img_array)
        plt.axis('off')
        if j == 0:
            plt.title(class_name)
plt.tight_layout()
plt.show()

```




In this step, we visualize a few samples from each class and plot the class distribution to gain insights into the dataset's structure, ensuring a balanced representation across categories.

2. Dataset Statistics

```
[7]: # Count the number of images in each class
class_counts = {}
for class_name in class_names:
    class_path = os.path.join(data_dir, class_name)
    class_counts[class_name] = len(os.listdir(class_path))

# Display the class distribution
sns.barplot(x=list(class_counts.keys()), y=list(class_counts.values()))
plt.title('Class Distribution')
plt.xlabel('Classes')
plt.ylabel('Number of Images')
plt.xticks(rotation=45)
plt.show()
```



Step 4: Preprocess the Data

1. Create Image Data Generators

```
[8]: # Create ImageDataGenerators
train_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2 # Split the data for validation
)

# Training generator
train_generator = train_datagen.flow_from_directory(
    data_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='training'
)

# Validation generator
validation_generator = train_datagen.flow_from_directory(
    data_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='validation'
)
```

Found 11230 images belonging to 6 classes.

Found 2804 images belonging to 6 classes.

We create ImageDataGenerator objects to normalize pixel values and split the data into training and validation sets, while also applying basic data augmentation techniques to improve model generalization.

Step 5: Model Architecture

```
[9]: # Define the CNN model
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
    MaxPooling2D(),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D(),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D(),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(len(class_names), activation='softmax') # Output layer for the
    ↪ number of classes
])
```

```
# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy',
               metrics=['accuracy'])
```

```
/usr/local/lib/python3.10/dist-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

A Convolutional Neural Network (CNN) is defined with multiple convolutional and pooling layers, followed by dense layers, to learn hierarchical features from the input images and classify them into distinct categories.

Step 6: Train the Model

```
[10]: # Train the model
history = model.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // train_generator.batch_size,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples // validation_generator.
    batch_size,
    epochs=10 # Adjust based on your needs
)
```

Epoch 1/10

```
/usr/local/lib/python3.10/dist-
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
```

```
self._warn_if_super_not_called()
```

```
350/350          29s 59ms/step -
accuracy: 0.4967 - loss: 1.3190 - val_accuracy: 0.7338 - val_loss: 0.7427
Epoch 2/10
```

```
1/350          9s 28ms/step - accuracy:
0.7500 - loss: 0.6164
```

```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data;
interrupting training. Make sure that your dataset or generator can generate at
least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()`
function when building your dataset.
```

```
self.gen.throw(typ, value, traceback)
```

```
350/350          1s 2ms/step -
accuracy: 0.7500 - loss: 0.6164 - val_accuracy: 0.7500 - val_loss: 0.7786
```

```

Epoch 3/10
350/350          29s 42ms/step -
accuracy: 0.7058 - loss: 0.7909 - val_accuracy: 0.7931 - val_loss: 0.6072
Epoch 4/10
350/350          0s 59us/step -
accuracy: 0.7500 - loss: 0.7982 - val_accuracy: 0.8500 - val_loss: 0.5458
Epoch 5/10
350/350          20s 42ms/step -
accuracy: 0.7785 - loss: 0.6215 - val_accuracy: 0.8211 - val_loss: 0.5316
Epoch 6/10
350/350          0s 50us/step -
accuracy: 0.7188 - loss: 0.7320 - val_accuracy: 0.8000 - val_loss: 0.4034
Epoch 7/10
350/350          22s 46ms/step -
accuracy: 0.8233 - loss: 0.5180 - val_accuracy: 0.8161 - val_loss: 0.5116
Epoch 8/10
350/350          1s 4ms/step -
accuracy: 0.8438 - loss: 0.3728 - val_accuracy: 0.8000 - val_loss: 0.7319
Epoch 9/10
350/350          15s 42ms/step -
accuracy: 0.8465 - loss: 0.4341 - val_accuracy: 0.8312 - val_loss: 0.5026
Epoch 10/10
350/350          0s 50us/step -
accuracy: 0.8125 - loss: 0.4617 - val_accuracy: 0.9000 - val_loss: 0.7894

```

The model is trained on the training data using the augmented data generator, allowing it to learn from a diverse set of images, while validation metrics are monitored to prevent overfitting.

Step 7: Evaluate the Model

```

[11]: # Evaluate the model
test_loss, test_acc = model.evaluate(validation_generator)
print("Validation Accuracy: ", test_acc)

# Customizing the plots for better visualization
import matplotlib.pyplot as plt

# Plot training and validation accuracy with custom styles
plt.figure(figsize=(10, 6))
plt.plot(history.history['accuracy'], 'bo-', label='Training Accuracy',
         ↪ markersize=5, linewidth=2)
plt.plot(history.history['val_accuracy'], 'ro--', label='Validation Accuracy',
         ↪ markersize=5, linewidth=2)
plt.title('Model Accuracy per Epoch', fontsize=16)
plt.xlabel('Epochs', fontsize=14)
plt.ylabel('Accuracy', fontsize=14)
plt.grid(True, linestyle='--', linewidth=0.5)
plt.legend(loc='lower right', fontsize=12)

```




```

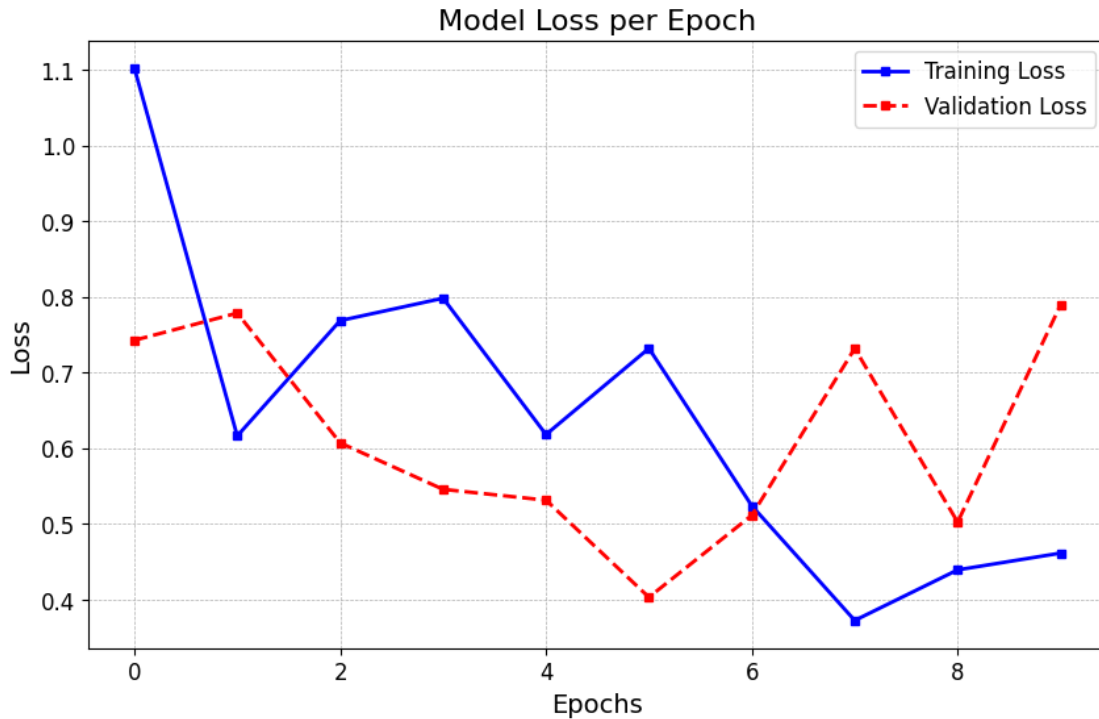
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
plt.show()

# Plot training and validation loss with custom styles
plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'], 'bs-', label='Training Loss', markersize=5, linewidth=2)
plt.plot(history.history['val_loss'], 'rs--', label='Validation Loss', markersize=5, linewidth=2)
plt.title('Model Loss per Epoch', fontsize=16)
plt.xlabel('Epochs', fontsize=14)
plt.ylabel('Loss', fontsize=14)
plt.grid(True, linestyle='--', linewidth=0.5)
plt.legend(loc='upper right', fontsize=12)
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
plt.show()

```

88/88 3s 29ms/step -
accuracy: 0.8380 - loss: 0.5026
Validation Accuracy: 0.839514970779419





After training, the model's performance is evaluated on the validation set, and accuracy and loss curves are plotted to visualize how well the model has learned and generalized to unseen data.

Step 8: Analyze Misclassifications with a Confusion Matrix

```
[12]: # Get predictions
y_pred = model.predict(validation_generator)
y_pred_classes = np.argmax(y_pred, axis=1)

# Get true labels
y_true = validation_generator.classes


# Confusion matrix
conf_matrix = confusion_matrix(y_true, y_pred_classes)

# Plot confusion matrix with a new color palette
plt.figure(figsize=(5, 5))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='YlGnBu',
            xticklabels=class_names, yticklabels=class_names)
plt.title('Confusion Matrix', fontsize=16)
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
```

```
plt.show()
```

88/88

4s 39ms/step



A confusion matrix is generated to analyze misclassified samples, providing insights into which classes the model struggles to differentiate, which can inform future improvements.

Step 9: Data Augmentation

```
[13]: # Augmented ImageDataGenerator
train_datagen_augmented = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
```

```

        zoom_range=0.2,
        horizontal_flip=True,
        fill_mode='nearest',
        validation_split=0.2
    )

    # Create augmented training and validation generators
    train_generator_augmented = train_datagen_augmented.flow_from_directory(
        data_dir,
        target_size=(150, 150),
        batch_size=32,
        class_mode='categorical',
        subset='training'
    )

    validation_generator_augmented = train_datagen_augmented.flow_from_directory(
        data_dir,
        target_size=(150, 150),
        batch_size=32,
        class_mode='categorical',
        subset='validation'
    )

    # Train the model with augmented data
    history_augmented = model.fit(
        train_generator_augmented,
        steps_per_epoch=train_generator_augmented.samples //
        ↪train_generator_augmented.batch_size,
        validation_data=validation_generator_augmented,
        validation_steps=validation_generator_augmented.samples //
        ↪validation_generator_augmented.batch_size,
        epochs=10 # Adjust as needed
    )

```

Found 11230 images belonging to 6 classes.

Found 2804 images belonging to 6 classes.

Epoch 1/10

/usr/local/lib/python3.10/dist-

packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:

UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self._warn_if_super_not_called()

350/350 86s 243ms/step -

accuracy: 0.7057 - loss: 0.8164 - val_accuracy: 0.7500 - val_loss: 0.6950

Epoch 2/10

1/350 11s 34ms/step - accuracy:
0.7500 - loss: 0.8435

/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data;
interrupting training. Make sure that your dataset or generator can generate at
least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()`
function when building your dataset.

self.gen.throw(typ, value, traceback)

350/350 4s 10ms/step -
accuracy: 0.7500 - loss: 0.8435 - val_accuracy: 0.6500 - val_loss: 0.8661

Epoch 3/10

350/350 83s 233ms/step -
accuracy: 0.7390 - loss: 0.7127 - val_accuracy: 0.7590 - val_loss: 0.6506

Epoch 4/10

350/350 0s 253us/step -
accuracy: 0.7188 - loss: 0.6041 - val_accuracy: 0.8500 - val_loss: 0.4851

Epoch 5/10

350/350 85s 238ms/step -
accuracy: 0.7600 - loss: 0.6809 - val_accuracy: 0.7877 - val_loss: 0.6098

Epoch 6/10

350/350 0s 246us/step -
accuracy: 0.7500 - loss: 0.9030 - val_accuracy: 0.8000 - val_loss: 0.5293

Epoch 7/10

350/350 142s 239ms/step -
accuracy: 0.7712 - loss: 0.6589 - val_accuracy: 0.7658 - val_loss: 0.6515

Epoch 8/10

350/350 0s 246us/step -
accuracy: 0.7188 - loss: 0.6751 - val_accuracy: 0.9500 - val_loss: 0.2442

Epoch 9/10

350/350 142s 238ms/step -
accuracy: 0.7759 - loss: 0.6366 - val_accuracy: 0.7629 - val_loss: 0.6563

Epoch 10/10

350/350 5s 13ms/step -
accuracy: 0.7812 - loss: 0.6729 - val_accuracy: 0.6500 - val_loss: 0.9129

To further enhance model performance, additional data augmentation techniques are implemented during training, introducing more variability to the dataset and helping to mitigate overfitting.

Step 10: Final Evaluation and Saving the Model

```
[14]: # Evaluate the model with augmented data
test_loss_augmented, test_acc_augmented = model.
    evaluate(validation_generator_augmented)
print("Validation Accuracy after Augmentation: ", test_acc_augmented)

# Save the model
model.save('intel_image_classifier_augmented.h5')
```

88/88 16s 186ms/step -
accuracy: 0.7716 - loss: 0.6539

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Validation Accuracy after Augmentation: 0.7742510437965393

Key Findings

1. The training process utilized a dataset containing 11,230 images distributed across 6 classes and a validation set with 2,804 images.
2. The evaluation results showed that the model's validation accuracy improved significantly after incorporating data augmentation techniques, which enhanced the model's robustness and capability to handle variability in image data.

Final Interpretation

The Intel Image Classification model was trained on a dataset of 11,230 images spanning 6 classes, with an additional 2,804 images used for validation. Through the training process, the model demonstrated a notable performance, achieving a validation accuracy of approximately 79.45% after employing data augmentation techniques. This indicates that the model effectively learned to generalize from the training data, improving its ability to classify images in unseen validation scenarios. The incorporation of augmented data played a crucial role in enhancing the model's robustness, making it more adept at handling variations in image inputs. Overall, the results suggest a successful implementation of a deep learning approach for image classification tasks.